

天池大赛

扶持计划

## 使用pg\_resetxlog修复PostgreSQL控制文件的方法

德哥 

2016-08-14

9884浏览量

**简介：** PostgreSQL 控制文件在\$PGDATA/global目录下名为pg\_control. 控制文件中记录了以下三部分信息：

1. initdb时生成的静态信息：pg\_control version number: 922 C

PostgreSQL 控制文件在\$PGDATA/global目录下名为pg\_control.

控制文件中记录了以下三部分信息：

1. initdb时生成的静态信息：

```
pg_control version number:          922
Catalog version number:            201204301
Database system identifier:         5831753892046499175
Maximum data alignment:             8
Database block size:                8192
Blocks per segment of large relation: 131072
WAL block size:                    16384
Bytes per WAL segment:              16777216
Maximum length of identifiers:       64
Maximum columns in an index:         32
Maximum size of a TOAST chunk:       1996
Date/time type storage:              64-bit integers
Float4 argument passing:             by value
Float8 argument passing:             by value
```

2. postgresql.conf中的配置信息：

```
Current wal_level setting:          hot_standby
Current max_connections setting:     1000
Current max_prepared_xacts setting:  10
Current max_locks_per_xact setting:  64
```

3. write-ahead logging以及checkpoint的动态信息：

```
Latest checkpoint location:          96E8/5B000020
Prior checkpoint location:            96E8/5A0C8CC0
Latest checkpoint's REDO location:    96E8/5B000020
Latest checkpoint's TimeLineID:       1
Latest checkpoint's full_page_writes: on
Latest checkpoint's NextXID:          0/1183664222
Latest checkpoint's NextOID:          309701
Latest checkpoint's NextMultiXactId:  1
Latest checkpoint's NextMultiOffset:  0
Latest checkpoint's oldestXID:        1006759584
Latest checkpoint's oldestXID's DB:    1
Latest checkpoint's oldestActiveXID:   0
Time of latest checkpoint:             Fri 11 Jan 2013 07:44:19 AM CST
Minimum recovery ending location:      0/0
```

```

Maximum recovery ending location: 0/0
Backup start location:             0/0
Backup end location:               0/0
End-of-backup record required:     no

```

以上信息可以使用pg\_controldata从pg\_control获取：

```

src/bin/pg_controldata/pg_controldata.c
* pg_controldata
*
* reads the data from $PGDATA/global/pg_control

```

如果控制文件\$PGDATA/global/pg\_control损坏或丢失, 数据库将运行异常, 无法启动.

如何修复? 关键在于恢复write-ahead logging以及checkpoint的动态信息.

这些信息可以从pg\_xlog, pg\_clog, pg\_multixact这些目录的文件中解析出来.

pg\_xlog的文件名解析可参看, 不同的段大小, 命名大不相同, pg\_resetxlog的帮助文件适用16MB的段大小, 如果是其他大小, 需要重新计算名字:

<http://blog.163.com/digoal@126/blog/static/1638770402012914112949546/>

接下来介绍一下使用pg\_resetxlog重建pg\_control的方法.

pg\_resetxlog功能如下:

```

src/bin/pg_resetxlog/pg_resetxlog.c
* pg_resetxlog.c
*
*   A utility to "zero out" the xlog when it's corrupt beyond recovery.
*   Can also rebuild pg_control if needed.
*
* The theory of operation is fairly simple:
*
*   1. Read the existing pg_control (which will include the last
*      checkpoint record). If it is an old format then update to
*      current format.
*
*   2. If pg_control is corrupt, attempt to intuit reasonable values,
*      by scanning the old xlog if necessary.
*
*   3. Modify pg_control to reflect a "shutdown" state with a checkpoint
*      record at the start of xlog.
*
*   4. Flush the existing xlog files and write a new segment with
*      just a checkpoint record in it. The new segment is positioned
*      just past the end of the old xlog, so that existing LSNs in
*      data pages will appear to be "in the past".
*
* This is all pretty straightforward except for the intuition part of
* step 2 ...

```



用法:

```

ocz@db-172-16-3-150-> pg_resetxlog --help
pg_resetxlog resets the PostgreSQL transaction log

```

pg\_resetxlog resets the PostgreSQL transaction log.

#### Usage:

pg\_resetxlog [OPTION]... DATADIR

#### Options:

-e XIDPOCH set next transaction ID epoch  
 -f force update to be done  
 -l TLI,FILE,SEG force minimum WAL starting location for new transaction log  
 -m XID set next multitransaction ID  
 -n no update, just show extracted control values (for testing)  
 -o OID set next OID  
 -O OFFSET set next multitransaction offset  
 -V, --version output version information, then exit  
 -x XID set next transaction ID  
 -?, --help show this help, then exit

#### 参数具体含义：

##### -l timelineid,fileid,seg

The WAL starting address (-l) should be larger than any WAL segment file name currently in the directory pg\_xlog under the data directory. These names are also in hexadecimal and have the first part is the "timeline ID" and should usually be kept the same. Do not choose a value (0xFF) for the third part; instead increment the second part and reset the third part to 0. 00000001000000320000004A is the largest entry in pg\_xlog, -l 0x1,0x32,0x4B will work; but if the entry is 000000010000003A000000FF, choose -l 0x1,0x3B,0x0 or more.

##### Note

pg\_resetxlog itself looks at the files in pg\_xlog and chooses a default -l setting beyond the existing file name. Therefore, manual adjustment of -l should only be needed if you are adding new segment files that are not currently present in pg\_xlog, such as entries in an offline backup. If the contents of pg\_xlog have been lost entirely.

##### -e XIDPOCH

The transaction ID epoch is not actually stored anywhere in the database except in the file pg\_resetxlog, so any value will work so far as the database itself is concerned. You might want to ensure that replication systems such as Slony-I work correctly - if so, an appropriate value should be obtainable from the state of the downstream replicated database.

##### -x XID

A safe value for the next transaction ID (-x) can be determined by looking for the numerical value in the directory pg\_clog under the data directory, adding one, and then multiplying by that the file names are in hexadecimal. It is usually easiest to specify the option value in hexadecimal. For example, if 0011 is the largest entry in pg\_clog, -x 0x120000 will work (five times the proper multiplier).

##### -m XID

A safe value for the next multitransaction ID (-m) can be determined by looking for the numerical value in the directory pg\_multixact/offsets under the data directory, adding one, and then multiplying by 65536. As above, the file names are in hexadecimal, so the easiest way to do this is to specify the value in hexadecimal and add four zeroes.

##### -O

A safe value for the next multitransaction offset (-O) can be determined by looking for

largest file name in the directory pg\_multixact/members under the data directory, adding on multiplying by 65536. As above, the file names are in hexadecimal, so the easiest way to do specify the option value in hexadecimal and add four zeroes.

-o OID

There is no comparably easy way to determine a next OID that's beyond the largest one in fortunately it is not critical to get the next-OID setting right.

-n

no update, just show extracted control values (for testing)

-f

force

测试步骤如下(基于PostgreSQL 9.2.1) :

1. 新建测试数据, 用到with oids的表, 因为OID无法确定, 看看是否有异常.
2. 关闭数据库
3. 记下pg\_controldata信息, 方便修复后进行比对
4. 删除\$PGDATA/global/pg\_control
5. 开启数据库观察报错输出
6. touch \$PGDATA/global/pg\_control
7. 使用pg\_resetxlog修复pg\_control
8. 记下pg\_controldata信息, 与前面的pg\_controldata输出进行比对
9. 启动数据库
10. 查看测试数据是否正常, 新插入数据
11. 关闭数据库, 并记下pg\_controldata的信息, 看看有何变化.

测试过程:

## 1. 测试数据

```

digoal=> create table oid_test(id int primary key) with oids;
NOTICE: CREATE TABLE / PRIMARY KEY will create implicit index "oid_test_pkey" for table "o
CREATE TABLE
digoal=> insert into oid_test select generate_series(1,100000);
INSERT 0 100000
digoal=> select min(oid),max(oid) from oid_test ;
   min   |   max
-----+-----
 16397 | 116396
(1 row)

```

## 2. 关闭数据

```
ocz@db-172-16-3-150-> pg_ctl stop -m fast
waiting for server to shut down.... done
server stopped
```

### 3. 记下pg\_controldata信息, 方便修复后进行比对

```
ocz@db-172-16-3-150-> pg_controldata
pg_control version number:          922
Catalog version number:            201204301
Database system identifier:         5832000131111550393
Database cluster state:             shut down
pg_control last modified:           Fri 11 Jan 2013 09:48:18 AM CST
Latest checkpoint location:         96E8/5F000020
Prior checkpoint location:          96E8/5EE5C698
Latest checkpoint's REDO location:  96E8/5F000020
Latest checkpoint's TimeLineID:     1
Latest checkpoint's full_page_writes: on
Latest checkpoint's NextXID:        0/1183842312
Latest checkpoint's NextOID:        116414
Latest checkpoint's NextMultiXactId: 65536
Latest checkpoint's NextMultiOffset: 65536
Latest checkpoint's oldestXID:      1006759584
Latest checkpoint's oldestXID's DB: 1
Latest checkpoint's oldestActiveXID: 0
Time of latest checkpoint:          Fri 11 Jan 2013 09:48:18 AM CST
Minimum recovery ending location:    0/0
Backup start location:               0/0
Backup end location:                 0/0
End-of-backup record required:       no
Current wal_level setting:           hot_standby
Current max_connections setting:     1000
Current max_prepared_xacts setting:  10
Current max_locks_per_xact setting:  64
Maximum data alignment:              8
Database block size:                 8192
Blocks per segment of large relation: 131072
WAL block size:                      16384
Bytes per WAL segment:               16777216
Maximum length of identifiers:       64
Maximum columns in an index:         32
Maximum size of a TOAST chunk:       1996
Date/time type storage:               64-bit integers
Float4 argument passing:              by value
Float8 argument passing:              by value
```

### 4. 删除\$PGDATA/global/pg\_control



```
ocz@db-172-16-3-150-> rm $PGDATA/global/pg_control
rm regular file `/data05/ocz/pg_root/global/pg_control'? y
```

## 5. 开启数据库观察报错输出

```
oczd@db-172-16-3-150-> pg_ctl start
server starting
oczd@db-172-16-3-150-> postgres: could not find the database system
Expected to find it in the directory "/data05/oczd/pg_root",
but could not open file "/data05/oczd/pg_root/global/pg_control": No such file or directory
```

接下来进行修复：

## 6. touch \$PGDATA/global/pg\_control

```
oczd@db-172-16-3-150-> touch $PGDATA/global/pg_control
oczd@db-172-16-3-150-> chmod 600 $PGDATA/global/pg_control
```

## 7. 使用pg\_resetxlog修复pg\_control

首先确定-l timelineid,fileid,seg的信息：

```
oczd@db-172-16-3-150-> cd $PGDATA/pg_xlog
oczd@db-172-16-3-150-> ll
total 65M
-rw----- 1 oczd oczd 16M Jan 11 09:39 00000001000096E80000005C
-rw----- 1 oczd oczd 16M Jan 11 09:39 00000001000096E80000005D
-rw----- 1 oczd oczd 16M Jan 11 09:48 00000001000096E80000005E
-rw----- 1 oczd oczd 16M Jan 11 09:48 00000001000096E80000005F
drwx----- 2 oczd oczd 44K Jan 11 09:48 archive_status
```

-l timelineid,fileid,seg 的数据来自pg\_xlog文件名的三个部分, 分别占用8个16进制位。

段大小为16MB, 所以末端最大为0xFF.

得出-l 0x1,0x96E8,0x60

接下来确定-x XID的信息

来自pg\_clog

```
oczd@db-172-16-3-150-> cd $PGDATA/pg_clog
oczd@db-172-16-3-150-> ll -t|head -n 5
total 43M
-rw----- 1 oczd oczd 8.0K Jan 11 09:48 0469
-rw----- 1 oczd oczd 216K Jan 10 21:00 0468
-rw----- 1 oczd oczd 256K Jan 10 12:56 0467
-rw----- 1 oczd oczd 256K Jan 10 09:35 0466
```

取最大值加1然后乘以1048576.

转换成16进制的话相当于取最大值加1然后末尾添加5个0

得到-x 0x046A00000

接下来确定-m XID的信息

来自pg\_multixact/offsets

```
ocz@db-172-16-3-150-> cd $PGDATA/pg_multixact/offsets
ocz@db-172-16-3-150-> ll
total 0
```

取最大值加1然后乘以65536.

转换成16进制的话相当于取最大值加1然后末尾添加4个0

没有文件的话使用0加1, 然后末尾添加4个0

得到-m 0x10000

接下来确定-O OFFSET的信息

来自pg\_multixact/members

```
ocz@db-172-16-3-150-> cd $PGDATA/pg_multixact/members
ocz@db-172-16-3-150-> ll
total 0
```

取最大值加1然后乘以65536.

转换成16进制的话相当于取最大值加1然后末尾添加4个0

没有文件的话使用0加1, 然后末尾添加4个0

得到-O 0x10000

最后, 不确定的值有2个:

-e XIDPOCH , 如果有使用 slony或者londiste这种基于触发器的数据同步软件, 则-e意义不大, 它实际  
-o OID , 系统会自动跳过已经分配的OID, 自动容错, 例如OID被别的程序使用掉了, PG会自动生成下一个OI



可以 5个值.

执行pg\_resetxlog 如下：

```
ocz@db-172-16-3-150-> pg_resetxlog -l 0x1,0x96E8,0x60 -x 0x046A00000 -m 0x10000 -O 0x10000
pg_resetxlog: pg_control exists but is broken or unknown version; ignoring it
Transaction log reset
```

8. 记下pg\_controldata信息, 与前面的pg\_controldata输出进行比对

```
ocz@db-172-16-3-150-> pg_controldata
pg_control version number:          922
Catalog version number:            201204301
Database system identifier:         5832008033851373032
Database cluster state:             shut down
pg_control last modified:           Fri 11 Jan 2013 10:09:44 AM CST
Latest checkpoint location:         96E8/60000020
Prior checkpoint location:          0/0
Latest checkpoint's REDO location:  96E8/60000020
Latest checkpoint's TimeLineID:     1
Latest checkpoint's full_page_writes: off
Latest checkpoint's NextXID:        0/1184890880
Latest checkpoint's NextOID:        10000
Latest checkpoint's NextMultiXactId: 65536
Latest checkpoint's NextMultiOffset: 65536
Latest checkpoint's oldestXID:      3479858176
Latest checkpoint's oldestXID's DB: 0
Latest checkpoint's oldestActiveXID: 0
Time of latest checkpoint:          Fri 11 Jan 2013 10:09:44 AM CST
Minimum recovery ending location:    0/0
Backup start location:               0/0
Backup end location:                 0/0
End-of-backup record required:       no
Current wal_level setting:           minimal
Current max_connections setting:     100
Current max_prepared_xacts setting:  0
Current max_locks_per_xact setting:  64
Maximum data alignment:              8
Database block size:                 8192
Blocks per segment of large relation: 131072
WAL block size:                      16384
Bytes per WAL segment:               16777216
Maximum length of identifiers:       64
Maximum columns in an index:         32
Maximum size of a TOAST chunk:       1996
Date/time type storage:               64-bit integers
Float4 argument passing:              by value
Float8 argument passing:              by value
```

注意 控制文件读取到的不确定的-e XIDPOCH和-o OID信息如下：





也就是initdb后的初始值.

```
Latest checkpoint's NextXID:      0/1184890880 : XIDPOCH=0
Latest checkpoint's NextOID:      10000
```

与修复pg\_control前发生了变化的值如下:

修复前

```
Database system identifier:      5832000131111550393
pg_control last modified:      Fri 11 Jan 2013 09:48:18 AM CST
Prior checkpoint location:      96E8/5EE5C698
Latest checkpoint's full_page_writes: on
Latest checkpoint's NextXID:      0/1183842312
Latest checkpoint's NextOID:      116414
Latest checkpoint's oldestXID:      1006759584
Latest checkpoint's oldestXID's DB: 1
Time of latest checkpoint:      Fri 11 Jan 2013 09:48:18 AM CST
Current wal_level setting:      hot_standby
Current max_connections setting: 1000
Current max_prepared_xacts setting: 10
```

修复后

```
Database system identifier:      5832008033851373032
pg_control last modified:      Fri 11 Jan 2013 10:09:44 AM CST
Prior checkpoint location:      0/0
Latest checkpoint's full_page_writes: off
Latest checkpoint's NextXID:      0/1184890880
Latest checkpoint's NextOID:      10000
Latest checkpoint's oldestXID:      3479858176
Latest checkpoint's oldestXID's DB: 0
Time of latest checkpoint:      Fri 11 Jan 2013 10:09:44 AM CST
Current wal_level setting:      minimal
Current max_connections setting: 100
Current max_prepared_xacts setting: 0
```

## 9. 启动数据库

```
ocz@db-172-16-3-150-> pg_ctl start
server starting
ocz@db-172-16-3-150-> LOG:  00000: loaded library "pg_stat_statements"
LOCATION:  load_libraries, miscinit.c:1249
```

## 10. 是否正常, 然后新插入数据

```
ocz@db-172-16-3-150-> psql digoal digoal
nsal (9.2.1)
```

```

Type "help" for help.
digoal=> select min(oid),max(oid),count(*) from oid_test ;
   min |   max   | count
-----+-----+-----
 16397 | 116396 | 100000
(1 row)

```

数据可以正常访问.

新插入数据:

```

digoal=> insert into oid_test select generate_series(100001,200000);
INSERT 0 100000
digoal=> select min(oid),max(oid),count(*) from oid_test ;
   min |   max   | count
-----+-----+-----
 16384 | 116396 | 200000
(1 row)
digoal=> select oid,* from oid_test where oid=16397;
   oid |   id
-----+-----
 16397 |     1
 16397 | 100014
(2 rows)

```

注意oid出现了重复, 印证了PostgreSQL中的说明, OID不确保唯一性.

11. 关闭数据库, 并记下pg\_controldata的信息, 看看有何变化.

```

ocz@db-172-16-3-150-> pg_ctl stop -m fast
waiting for server to shut down.... done
server stopped
ocz@db-172-16-3-150-> pg_controldata
pg_control version number:          922
Catalog version number:            201204301
Database system identifier:         5832008033851373032
Database cluster state:             shut down
pg_control last modified:           Fri 11 Jan 2013 10:16:18 AM CST
Latest checkpoint location:         96E8/61000020
Prior checkpoint location:          96E8/60DFF470
Latest checkpoint's REDO location:  96E8/61000020
Latest checkpoint's TimeLineID:     1
Latest checkpoint's full_page_writes: on
Latest checkpoint's NextXID:        0/1184890883
Latest checkpoint's NextOID:        116385
Latest checkpoint's NextMultiXactId: 65536
Latest checkpoint's NextMultiOffset: 65536
Latest checkpoint's oldestXID:      1006759584
Latest checkpoint's oldestXID's DB: 1
Latest checkpoint's oldestActiveXID: 0
Time of latest checkpoint:          Fri 11 Jan 2013 10:16:18 AM CST

```

```

-----
Minimum recovery ending location:      0/0
Backup start location:                 0/0
Backup end location:                   0/0
End-of-backup record required:         no
Current wal_level setting:             hot_standby
Current max_connections setting:       1000
Current max_prepared_xacts setting:    10
Current max_locks_per_xact setting:    64
Maximum data alignment:                8
Database block size:                  8192
Blocks per segment of large relation: 131072
WAL block size:                       16384
Bytes per WAL segment:                 16777216
Maximum length of identifiers:         64
Maximum columns in an index:           32
Maximum size of a TOAST chunk:         1996
Date/time type storage:                64-bit integers
Float4 argument passing:               by value
Float8 argument passing:               by value

```

关闭数据库后与刚修复好时的控制文件信息变化如下：

开库前：

```

pg_control last modified:      Fri 11 Jan 2013 10:09:44 AM CST
Latest checkpoint location:    96E8/60000020
Prior checkpoint location:     0/0
Latest checkpoint's REDO location: 96E8/60000020
Latest checkpoint's full_page_writes: off
Latest checkpoint's NextXID:   0/1184890880
Latest checkpoint's NextOID:   10000
Latest checkpoint's oldestXID: 3479858176
Latest checkpoint's oldestXID's DB: 0
Time of latest checkpoint:     Fri 11 Jan 2013 10:09:44 AM CST
Current wal_level setting:     minimal
Current max_connections setting: 100
Current max_prepared_xacts setting: 0

```

关库后：

```

pg_control last modified:      Fri 11 Jan 2013 10:16:18 AM CST
Latest checkpoint location:    96E8/61000020
Prior checkpoint location:     96E8/60DFF470
Latest checkpoint's REDO location: 96E8/61000020
Latest checkpoint's full_page_writes: on
Latest checkpoint's NextXID:   0/1184890883
Latest checkpoint's NextOID:   116385
Latest checkpoint's oldestXID: 1006759584
Latest checkpoint's oldestXID's DB: 1
Time of latest checkpoint:     Fri 11 Jan 2013 10:16:18 AM CST
Current wal_level setting:     hot_standby
Current max_connections setting: 1000

```

Current max\_prepared\_xacts setting: 10

## 小结

1. 使用pg\_resetxlog后, 先检查数据一致性, 必要时将数据导出, 使用initdb新建数据库, 再导入.
2. 如果控制文件丢失, 并且没有备份的话, pg\_resetxlog你不知道该填啥, 但是可以从pg\_xlog目录中获得大概的redo location, 或者pg\_resetxlog 会猜测一些值, 直接-f生成控制文件, 启动数据库后, 可能由于XID回归到以前的XID而致使数据"消失", 你可以使用txid\_current()函数不断的消耗XID来得到一致的值.

<http://blog.163.com/digoal@126/blog/static/163877040201183043153622/>

<http://blog.163.com/digoal@126/blog/static/163877040201251911813661/>

使用pg\_xlogdump从xlog中抽取信息, 包括txid.

<http://blog.163.com/digoal@126/blog/static/16387704020134993845555/>

pg\_resetxlog的版本必须要与数据库集群的版本一致。

## 注意

控制文件和PostgreSQL的大版本相关, 建议用户使用pg\_resetxlog前, 看一看对应版本的使用说明。

## 参考

1. man pg\_controldata
2. man pg\_resetxlog
3. <http://blog.163.com/digoal@126/blog/static/163877040201171233710582/>
4. <http://blog.163.com/digoal@126/blog/static/1638770402012914112949546/>
5. <http://www.postgresql.org/docs/9.2/static/wal-internals.html>
6. <http://www.postgresql.org/docs/9.2/static/app-pgresetxlog.html>
7. <http://www.postgresql.org/docs/9.2/static/pgupgrade.html>
8. <http://www.postgresql.org/docs/9.2/static/continuous-archiving.html>
9. src/bin/pg\_resetxlog/pg\_resetxlog.c
10. src/include/catalog/pg\_control.h
11. src/bin/pg\_controldata/pg\_controldata.c
12. src/backend/access/transam/clog.c
13. src/include/access/clog.h
14. src/backend/access/transam/xlog.c
15. src/include/access/xlog.h

祝大家工作顺利, 欢迎随时来 阿里云促膝长谈业务需求 , 恭候光临。

阿里云的小伙伴们加油，努力 **做好内核与服务，打造最贴地气的云数据库**。

推荐关注：

- 手慢无！云服务器免费领取，实时计算净省3258元
- 一键收藏！2020年度最受欢迎电子书出炉
- 每周二看直播，学最新阿里开源技术
- 学它！阿里专家领学，大数据必看课程

**版权声明：**本文内容由阿里云实名注册用户自发贡献，版权归原作者所有，阿里云开发者社区不拥有其著作权，亦不承担相应法律责任。具体规则请查看《[阿里云开发者社区用户服务协议](#)》和《[阿里云开发者社区知识产权保护指引](#)》。如果您发现本社区中有涉嫌抄袭的内容，填写[侵权投诉表单](#)进行举报，一经查实，本社区将立刻删除涉嫌侵权内容。

关系型数据库

数据库

PostgreSQL

## 评论

登录后可评论



pgfans

2016-08-16

德哥，我在9.4.5上重复同样的测试遇到一个问题就是 - O 参数的计算。现在的版本说明是：A safe value for the next multitransaction offset (-O) ===== can be determined by looking for the numerically largest file name in the directory pg\_multixact/members under the data directory, adding one, and then multiplying by 52352. As above, the file names are in hexadecimal. There is no simple recipe such as the ones above of appending zeroes. 如果我的数据库 pg\_multixact/members 的最后文件名是：40BF 那么 -O 这个参数应该是多少啊？

👍 5

pis-j

2017-03-28

德哥，pg\_control丢失倒是不太可能，还有什么其他场景需要用pg\_resetxlog去做数据库修复的吗？

👍

德哥

2016-08-17

你是什么版本的PG，就看对应版本的说明。不同的版本可能不一样。

pgfans

2016-08-17

我设置为-O 0x40C00000 重新创建pg\_control之后报错：could not access status of transaction 1086324736", "Could not open file ""pg\_multixact/members/510E"": No such file or directory.



pgfans

2016-08-17

是40C0\*52352 还是40C0\*65536。文档描述multiplying by 52352



德哥

2016-08-17

用对应版本的说明，9.4：40C0 \* 65536



## 阿里云数据库

[+ 订阅](#)

帮用户承担一切数据库风险，给您何止是安心！

### 官方博客

如宝马3系和5系：PolarDB-X 与 DRDS 并驾齐驱

PolarDB-X 存储架构之“基于Paxos的最佳生产实践”

让SQL优化器更准确！AnalyticDB PG版发布统计信息自动收集功能

PolarDB-X 私有协议：提升集群的性能和稳定性

今天10点，阿里云重磅出品《超级探访·春节特辑》首次上线

[展开](#)

### 链接

RDS-MySQL版: <https://cn.aliyun.com/product/rds>

MongoDB版: <https://cn.aliyun.com/product/mongodb>

HBase版: <https://cn.aliyun.com/product/hbase>

**推荐**



售前咨询	支持与服务	账户管理	快速入口	资源和社区	关注阿里云
95187转1	联系我们	管理控制台	域名信息查询	开发者社区	企业决策
专业技术咨询	支持计划	备案管理	( WHOIS )	开发者学院	必读
全方位产品解读	极速工单	域名控制台	云产品快报	公开课	阿里云
成熟解决方案	新手学堂	账号管理	海外上云	问答	APP
成功客户案例分享	帮助文档	充值付款	客户案例	藏经阁	阿里云微
	自助工具	线下汇款/电汇	举报中心	天池平台	信
	迁移与部署	合同申请	信任中心	阿里云创新中心	阿里云微
	运维与管理	索取发票	产品图标	阿里云大学	博
	优化与提升	合规性	产品学习路径	在线学习	阿里云支
	公告		阿里云认证	开放实验室	持与服务
	聆听·建议反馈		贡献者榜单	考试认证	
	服务案例			高校合作	
	专家分享			码上公益	
	上云能力中心				
	CMC				

热门产品	数据传输服	弹性容器实
	务	例ECI
	服务网络	配置审计
	多媒体AI	
	访问控制	
	运维编排	
	视频监控	
	智能外呼	
用户热搜	云数据库	
	RDS	
	VPN网关	
	服务	
	表格存储	
	互	
	堡垒机	
	prometheus	
	云呼叫中心	
	数据洞察	
企业场景	创新创业直	
	企业专场销	
	企业上云	
	企业建站	
	企业营销	
	播间	
	量榜	
	企业邮箱	
	企业安全	
	飞天会员	
	企业数据	
更多推荐		
	移动安全加	
	全站加速	
	风险识别	
	DDoS防护	
	专有云	
	安全管家	
	视频审核	
	块存储	
	短信服务	
	固	

关于阿里云2.0

法律声明及隐私权政策

廉正举报

联系我们

加入阿里云

阿里巴巴集团

淘宝网

天猫

聚划算

全球速卖通

阿里巴巴国际交易市场

1688

阿里妈妈

飞猪

阿里云计算

AliOS

阿里通信

万网

高德

UC

友盟

虾米

优酷

钉钉

支付宝

达摩院

阿里安全

淘宝海外

© 2009-2021 Aliyun.com 版权所有 增值电信业务经营许可证：浙B2-20080101

☐

☐

安备 33010602009975号 浙B2-20080101-4