



PostgreSQL 中文社区
www.postgres.cn



世界上功能最强大的
开源数据库

首页
了解PostgreSQL
社区新闻
软件下载
文档资料
商业支持
图书
关于中文社区
有问必答(FAQ)
V2
登录

- QQ群1: 5276420
- QQ群2: 3336901
- QQ群3: 254622631
- 文档群: 150657323
- 文档翻译平台: [按此](#)
- 邮件列表: [按此订](#)
[阅](#)

[首页](#) --> [社区新闻](#) --> [强人随笔](#)

PostgreSQL流复制案例分析 | Startup 进程waiting问题分析

原作者: 刘雪柏 创作时间: 2019-08-12 10:41:13+08

wangliyun 发布于2019-08-13 08:08:13

评论: 1

浏览: 2418

顶: 176

踩: 16

一、作者介绍

刘雪柏, 目前就职于诺基亚东软通信有限公司, 高级软件工程师, 从事数据库相关工作, 联系方式: aiden.liu@qq.com。

二、问题现象

Postgresql 11.2版本物理复制, startup 进程命令行有时会出现waiting 标识。本文分析了出现waiting 标识的原因。

三、问题分析:

通过ps -ef | grep startup 监控startup 进程。

postgres: startup process recovering 0000000100000000000000A6

当startup 进程出现waiting的标识。通过pstack查看startup 进程调用堆栈

[postgres@sscloud21 ~]\$ pstack 197120

#0 0x00007f695b6faf53 in _selectnocancel () from /lib64/libc.so.6

#1 0x000000000088c97a in pg_usleep (microsec=) at pgsleep.c:56

#2 0x0000000000729ef9 in WaitExceedsMaxStandbyDelay () at standby.c:201

#3 ResolveRecoveryConflictWithVirtualXIDs (waitlist=0x1c706e0, reason=reason@entry=PROCSIGRECOVERYCONFLICT_SNAPSHOT, w at standby.c:262

#4 0x000000000072a10e in ResolveRecoveryConflictWithVirtualXIDs (reason=PROCSIGRECOVERYCONFLICT_SNAPSHOT, w at standby.c:315

#5 ResolveRecoveryConflictWithSnapshot (latestRemovedXid=, node=...) at standby.c:313

#6 0x00000000004c23be in heapxlogclean (record=0x1c00698) at heapam.c:8198

#7 heap2_redo (record=0x1c00698) at heapam.c:9351

#8 0x0000000000503e85 in StartupXLOG () at xlog.c:7306

#9 0x00000000006d82b1 in StartupProcessMain () at startup.c:211

#10 0x0000000000512275 in AuxiliaryProcessMain (argc=argc@entry=2, argv=argv@entry=0x7fff8b5d99b0) at bootstrap.c:441

#11 0x00000000006d53a0 in StartChildProcess (type=StartupProcess) at postmaster.c:5331

#12 0x00000000006d7b75 in PostmasterMain (argc=argc@entry=3, argv=argv@entry=0x1bd0e40) at postmaster.c:1371

#13 0x000000000048124f in main (argc=3, argv=0x1bd0e40) at main.c:228

检查源码发现, startup 进程正在执行 XLOGHEAP2CLEAN 类型的redo操作 (vacuum命令产生的块清除redo)。

heapam.c代码中有如下注释:



官方微信公众号



官方微博

搜索:

用户:

密码:

[注册](#)

[晋级](#)

[PostgreSQL的企业级云
服务](#)



PostgreSQL携手阿里云



```

8189 /*
8190  * We're about to remove tuples. In Hot Standby mode, ensure that there'
8191  * no queries running for which the removed tuples are still visible.
8192  *
8193  * Not all HEAP2_CLEAN records remove tuples with xids, so we only want
8194  * conflict on the records that cause MVCC failures for user queries. If
8195  * latestRemovedXid is invalid, skip conflict processing.
8196  */
8197 if (InHotStandby && TransactionIdIsValid(xlrec->latestRemovedXid))
8198     ResolveRecoveryConflictWithSnapshot(xlrec->latestRemovedXid, rnode);

```

startup 进程在执行清除tuples前需要确认没有活动的查询, 并且这些tuples对于活动查询仍然可见。

跟着ResolveRecoveryConflictWithSnapshot->ResolveRecoveryConflictWithVirtualXIDs->WaitExceedsMaxStandbyDelay->GetStandbyLimitTime

standby.c代码中GetStandbyLimitTime函数:

```

148 /*
149  * Determine the cutoff time at which we want to start canceling conflicting
150  * transactions. Returns zero (a time safely in the past) if we are willing
151  * to wait forever.
152  */
153 static TimestampTz
154 GetStandbyLimitTime(void)
155 {
156     TimestampTz rtime;
157     bool        fromStream;
158
159     /*
160      * The cutoff time is the last WAL data receipt time plus the appropriate
161      * delay variable. Delay of -1 means wait forever.
162      */
163     GetXLogReceiptTime(&rtime, &fromStream);
164     if (fromStream)
165     {
166         if (max_standby_streaming_delay < 0)
167             return 0; /* wait forever */
168         return TimestampTzPlusMilliseconds(rtime, max_standby_streaming_delay);
169     }
170     else
171     {
172         if (max_standby_archive_delay < 0)
173             return 0; /* wait forever */
174         return TimestampTzPlusMilliseconds(rtime, max_standby_archive_delay);
175     }
176 }

```

从代码中可以看出startup 进程在执行清除tuples前需要确认没有活动的查询有等待时间限制, fromStream的情况超过

maxstandbystreamingdelay参数时间限制会kill掉正在执行的查询, 其他情况如果超过了maxstandbyarchivedelay

参数时间限制会调用CancelVirtualTransaction函数 kill掉正在执行的查询。

standby.c相关代码

```

261         /* Is it time to kill it? */
262         if (WaitExceedsMaxStandbyDelay())
263         {
264             pid t      pid;
265
266             /*
267              * Now find out who to throw out of the balloon.
268              */
269             Assert(VirtualTransactionIdIsValid(*waitlist));
270             pid = CancelVirtualTransaction(*waitlist, reason);
271
272             /*
273              * Wait a little bit for it to die so that we avoid flooding
274              * an unresponsive backend when system is heavily loaded.
275              */
276             if (pid != 0)
277                 pg_usleep(5000L);
278         }
279     }

```

procarray.c相关代码

```

2661 pid t
2662 CancelVirtualTransaction(VirtualTransactionId vxid, ProcSignalReason sigmoc
2663 {
2664     ProcArrayStruct *arrayP = procArray;
2665     int index;
2666     pid t      pid = 0;
2667
2668     LWLockAcquire(ProcArrayLock, LW_SHARED);
2669
2670     for (index = 0; index < arrayP->numProcs; index++)
2671     {
2672         int pgprocno = arrayP->pgprocnos[index];
2673         volatile PGPROC *proc = &allProcs[pgprocno];
2674         VirtualTransactionId procvxid;
2675
2676         GET_VXID_FROM_PGPROC(procvxid, *proc);
2677
2678         if (procvxid.backendId == vxid.backendId &&
2679             procvxid.localTransactionId == vxid.localTransactionId)
2680         {
2681             proc->recoveryConflictPending = true;
2682             pid = proc->pid;
2683             if (pid != 0)
2684             {
2685                 /*
2686                  * Kill the pid if it's still here. If not, that's what we
2687                  * wanted so ignore any errors.
2688                  */
2689                 (void) SendProcSignal(pid, sigmode, vxid.backendId);
2690             }
2691             break;
2692         }
2693     }
2694 }

```

四、小结

- 1、原因: 重放XLOGHEAP2CLEAN 类型的redo操作不允许有查询, 如果有会等待, 等待时间由maxstandbystreamingdelay/maxstandbyarchivedelay控制。
- 2、解决这个问题可以控制slave端禁止执行查询操作, 或者调整maxstandbystreamingdelay/maxstandbyarchivedelay参数的值到范围。
- 3、maxstandbystreamingdelay/maxstandbyarchivedelay参数默认值30秒。



评论:1 浏览: 2418

👍 顶: 176 👎 踩: 162

评论：请在[登录](#)后发表评论，否则无法保存。

1# [__hemuyang](#) 回答于 2019-09-15 10:54:23+08
请问下如何复现该问题

发表评论：

 [您还没有登录,请您登录后再发表评论](#)